

**Course: C553**  
**Cloud Computing**  
**PA 1**  
**Benchmarking**  
**Design Document and Performance**

**Name: Shushupti Vijay Ajmire**  
**Prasanna Shanmuganathan**

# **Design Document**

This project aims at Benchmarking using KVM virtual machine m1.medium (2 virtual processors with 4GB RAM and 40GB disk), CentOS 7 Linux for the OS.

## **CPU Benchmarking**

- For CPU Benchmarking, we used strong scaling experiment and design a source code in C language.
- This program will measure the processor speed, in terms of double precision floating point operations per second (GFLOPS) and integer operations per second (GIOPS).
- This program has 16 instruction which include addition operation on both integer and floating point numbers which will run for 1 billion times to measure the speed.
- Main aim of this is to utilize the complete CPU cycle by performing diverse types of instructions and measure the speed.
- There is a performance evaluation which evaluates IOPS and FLOPS executed per second using different levels of concurrency (1 thread, 2 threads, 4 threads, 8 threads)
- Also run this code with AVX instruction and adequate evaluation using bare metal provisioning.
- Ran the Linpack benchmark (<http://en.wikipedia.org/wiki/LINPACK>) and compare the performance achieved using double precision floating point.

## **Memory Benchmarking**

- For Memory Benchmarking, we used strong scaling experiment and design a source code in C language.
- This program will measure the memory speed of host which will include parameters like read + write operations, sequential write access, random write access.
- This experiment used different block sizes (8B, 8KB, 8MB, 80MB), and varying the concurrency (1 thread, 2 threads, 4 threads, and 8 threads).
- We allocate 1 GB of memory to perform read + write operations in both sequential and random write access.
- The program also measures the throughput (Megabytes per second, MB/sec) and latency (microseconds, us). 8B block case used to measure latency, and the 8KB, 8MB, and 80MB cases used to measure throughput.
- Computed the theoretical memory bandwidth of your memory, based on the type of memory and the speed.
- Ran the Stream benchmark (<http://www.cs.virginia.edu/stream/>) and compared both the performance to see weather benchmarking is running properly.

## **Disk Benchmarking**

- For Disk Benchmarking, we used strong scaling experiment and design a source code in C language.
- This program will measure the memory speed of host which will include parameters like read + write operations, sequential write access, random write access.
- This experiment used different block sizes (8B, 8KB, 8MB, 80MB), and varying the concurrency (1 thread, 2 threads, 4 threads, and 8 threads).
- It implements 4 methods, sequential read and write, random read and write.
- Allocated a large piece file of 1GB, and perform read + write or read operations on either sequential or random access within this 1GB file.
- Ran the IOZone benchmark (<http://www.iozone.org/>) and compared the performance achieved with theoretical performance.

## **Network Benchmarking**

- For Network Benchmarking, we used strong scaling experiment and design a source code in Python language.
- It will measure the network speed over the loopback interface card with 1 node, between 2 processes on the same node and 2 nodes, between 2 processes on the different nodes.
- It includes the TCP protocol stack, UDP, fixed packet/buffer size (64KB), and varying the concurrency (1 thread, 2 threads, 4 threads, 8 threads) on both server and client side.
- It will measure the throughput (Megabits per second, Mb/sec) and latency (ms).
- We used fixed amount of data (1 GB) to transfer using socket programming and multithreading.
- Ran the IPerf benchmark (<http://en.wikipedia.org/wiki/Iperf>) and compared achieved throughput performance over the loopback interface.
- Compared efficiency with the theoretical memory performance.
- Repeated the benchmark across two nodes and compared the efficiency with theoretical network speed.

# Performance

## CPU Benchmarking

Performance is done on Chameleon testbud using KVM Virtual machine i.e m1.medium(2 virtual processors with 4GB RAM and 40GB disk.

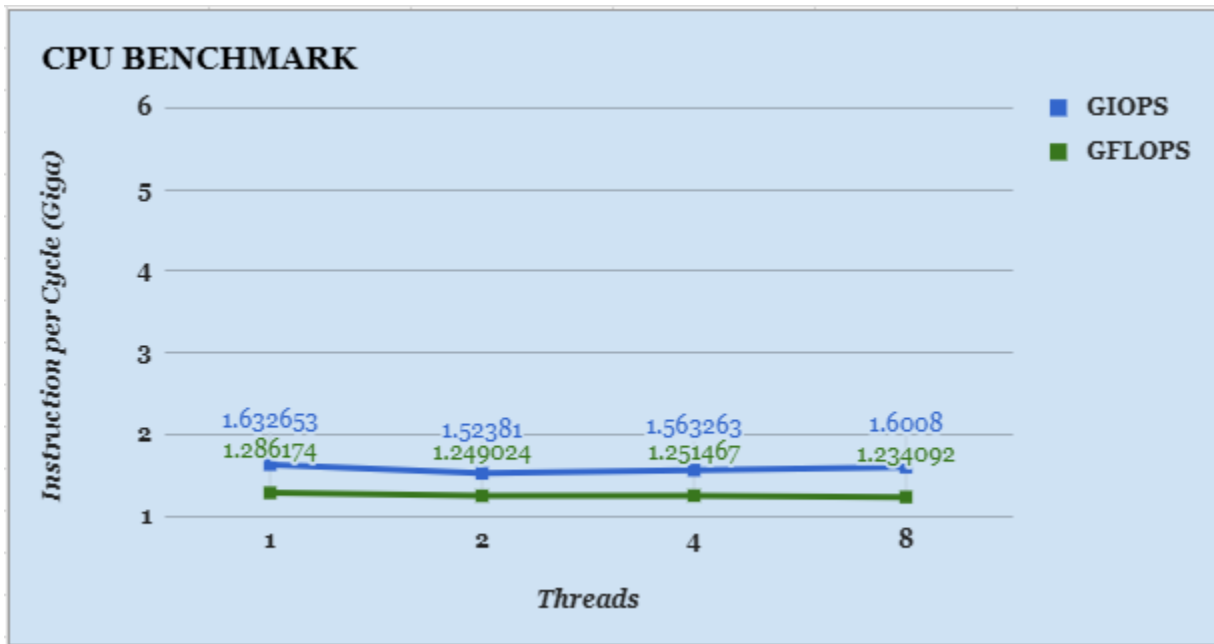
### Specifications:

Flavor	m1.medium
Flavor ID	3
RAM	4GB
VCPUs	2 VCPU
Disk	40GB

In this experiment, we are benchmarking CPU performance in terms of GIOPS and GFLOPS. And the experiment is done across 1,2,4 and 8 threads.

Below graphs shows the performance of the experiments:

1. This graph shows the Average GIOPS(Integer operations) and GFLOPS(Floating point operations) for 1,2,4,8 threads :



The above graph shows that the instructions per cycle decreasing after 1 thread and getting stabilized for the rest. It cannot be proven that if the thread count increases, the performance increases as well which can also

been seen above where In case of 1 threads, both GFLOPS and GIOPS have a higher value and then maintaining similar performance for 2,4,8 threads respectively.

### Test Results:

	GIOPS				GLOPS			
	1 thread	2 thread	4 thread	8 thread	1 thread	2 thread	4 thread	8 thread
Exp 1	1.632653	1.52381	1.563263	1.6008	1.286174	1.249024	1.251467	1.234092
Exp 2	1.596806	1.585728	1.59355	1.606023	1.244168	1.273885	1.24805	1.234806
Exp 3	1.62783	1.59542	1.602589	1.619758	1.269841	1.24031	1.273885	1.237424
Average	1.619096333	1.568319333	1.586467333	1.608860333	1.266727667	1.254406333	1.257800667	1.235440667
Std Dev	0.01588415463	0.03172059564	0.01681778214	0.0079953856	0.01728960463	0.01422550035	0.01145857194	0.001432400627

### Theoretical Performance

FLOPS= Number of Cores \* CLOCK SPEED \* IPC

$$= 2 * 2299.996 * 10^6 * 8 = 35.93 \text{ GFLOPS}$$

### Linpack Benchmark:

We ran the LINPACK Benchhmark on the system to compare the CPU performance and from the below it can be seen that the maximal value is obtained at 41.0406 GFLOPS which is 95% matching with the theoretical performance

```
[cc@pal-shushupti-prasanna linpack]$ cat lin_xeon64.txt
Sun Oct 8 20:46:03 UTC 2017
Intel(R) LINPACK data
Current date/time: Sun Oct 8 20:46:03 2017

CPU frequency: 3.095 GHz
Number of CPUs: 2
Number of cores: 2
Number of threads: 2

Parameters are set to:
Number of tests : 15
Number of equations to solve (problem size) : 1000 2000 5000 10000 15000 18000 20000 22000 25000 26000 27000 30000 35000 40000 45000
Leading dimension of array : 1000 2000 5000 10000 15000 18000 20016 22006 25000 26000 27000 30000 35000 40000 45000
Number of trials to run : 4 2 2 2 2 2 2 2 2 2 1 1 1 1 1
Data alignment value (in kbytes) : 4 4 4 4 4 4 4 4 4 4 1 1 1 1 1
Maximum memory requested that can be used = 3202964416, at the size = 20000

===== Timing linear equation system solver =====
Size LDA Align. Time(s) GFlops Residual Residual(norm)
1000 1000 4 0.022 30.6062 1.029343e-12 3.510325e-02
1000 1000 4 0.022 30.8068 1.029343e-12 3.510325e-02
1000 1000 4 0.023 28.7229 1.029343e-12 3.510325e-02
1000 1000 4 0.030 22.2561 1.029343e-12 3.510325e-02
2000 2000 4 0.153 34.8120 4.298950e-12 3.739560e-02
2000 2000 4 0.151 35.3859 4.298950e-12 3.739560e-02
5000 5000 4 2.157 36.6552 2.561643e-11 3.599893e-02
5000 5000 4 2.073 40.2211 2.561643e-11 3.599893e-02
10000 10000 4 16.606 39.6794 9.603002e-11 3.386116e-02
10000 10000 4 17.056 39.0989 9.603002e-11 3.386116e-02
15000 15000 4 55.769 40.3531 2.042799e-10 3.217442e-02
15000 15000 4 58.388 39.9104 2.042799e-10 3.217442e-02
18000 18000 4 97.344 39.9473 2.894967e-10 3.170367e-02
18000 18000 4 96.555 40.2741 2.894967e-10 3.170367e-02
20000 20016 4 129.972 41.0406 4.097966e-10 3.627616e-02
20000 20016 4 130.719 40.6061 4.097966e-10 3.627616e-02


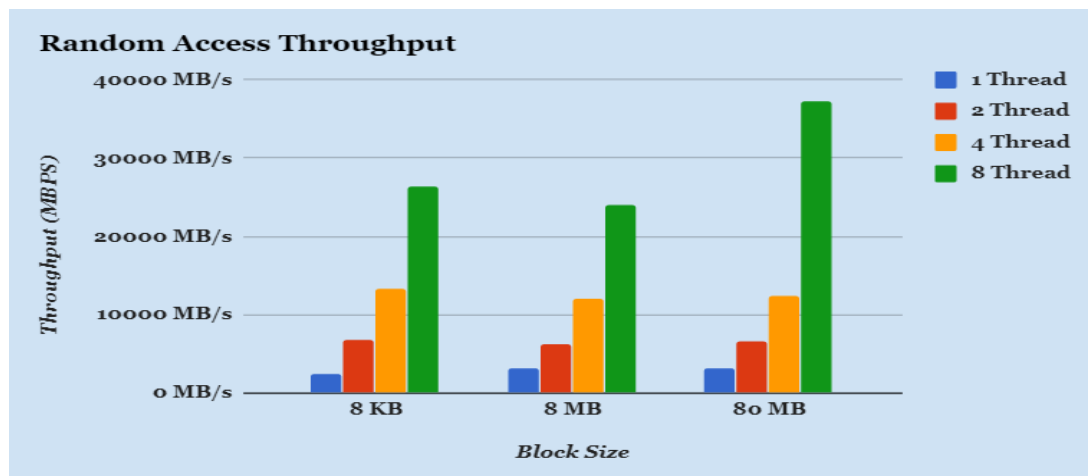

Performance Summary (GFlops)
Size LDA Align. Average Maximal
1000 1000 4 28.0985 30.8068
2000 2000 4 35.0989 35.3859
5000 5000 4 39.4381 40.2211
10000 10000 4 39.3891 39.6794
15000 15000 4 40.1317 40.3531
18000 18000 4 40.1107 40.2741
20000 20016 4 40.9233 41.0406


End of tests
Sun Oct 8 20:57:54 UTC 2017
[cc@pal-shushupti-prasanna linpack]$
```

## Memory Benchmarking

In this experiment, we are performing Memory Benchmarking across 1,2,4 and 8 threads to measure the efficiency of concurrency. Thread operations are done on different block size of data - 8B,8KB,8MB and 80MB and the operations are Sequential and random read+write.

Below are the graphs showing the experimental results of the program we written on C:





Throughputs were measured for 8kb,8MB and 80MB cases whereas latency was plotted for 8B block-size case. As we see from the above graphs, we could see that the throughput increasing as we increase the number of the threads. And the Latency is also decreasing on the thread increase. With sequential read access, the latency has seen a gradual decrease. However, with random access there was increase in latency upon using 2 threads to perform the read+write operation.

### Test Results :

#### Throughput:

We took 3 experiment results for statistics :

Test Results for 8B								
	Sequential R+W				Random R+W			
	1 Thread	2Thread	4Thread	8Thread	1 Thread	2Thread	4Thread	8Thread
Exp 1	1219.047619	2301.123596	4708.045977	9525.581395	202.371542	26.933193	50.44335	97.931859
Exp 2	1000	1684.210526	3368.421053	7111.111111	216.216216	26.845638	89.761571	92.352092
Exp 3	1000	1777.777778	3368.421053	6736.842105	219.178082	27.610009	53.333333	100.946372
Average	1073.015873	1921.0373	3814.962694	7791.178204	212.5886133	27.12961333	64.51275133	97.07677433
Standard Dev.	103.2600379	271.4625842	631.5052453	1235.889732	7.32505165	0.3415664571	17.89255285	3.560317357

Test Results for 8 KB								
	Sequential R+W				Random R+W			
	1 Thread	2Thread	4Thread	8Thread	1 Thread	2Thread	4Thread	8Thread
Exp 1	3200	6206.060606	12800	24824.24242	2438.095238	6826.666667	13212.90323	26425.80645
Exp 2	3200	6400	10666.66667	21333.33333	2666.666667	4571.428571	10666.66667	21333.33333
Exp 3	3200	6400	10666.66667	24824.24242	3200	4571.428571	10666.66667	21333.33333
Average	3200	6335.353535	11377.77778	23660.60606	2768.253968	5323.174603	11515.41219	23030.82437
Standard Dev.	0	91.42390709	1005.662976	1645.630327	319.2331451	1063.129434	1200.307425	2400.614851

Test Results for 8 MB								
	Sequential R+W				Random R+W			
	1 Thread	2Thread	4Thread	8Thread	1 Thread	2Thread	4Thread	8Thread
Exp 1	3011.764706	5688.888889	10778.94737	22140.54054	3103.030303	6206.060606	12047.05882	24094.11765
Exp 2	2666.666667	4571.428571	9142.857143	21333.33333	2666.666667	5333.333333	10666.66667	25600
Exp 3	2666.666667	5333.333333	10666.66667	25600	3200	4571.428571	9142.857143	21333.33333
Average	2781.699347	5197.883598	10196.15706	23024.62462	2989.89899	5370.27417	10618.86088	23675.81699
Standard Dev.	162.6807757	466.146875	746.2047426	1850.642162	231.9628296	667.8467596	1186.117164	1766.794215

Test Results for 80 MB								
	Sequential R+W				Random R+W			
	1 Thread	2Thread	4Thread	8Thread	1 Thread	2Thread	4Thread	8Thread
Exp 1	3413.333333	5851.428571	11702.85714	35617.3913	3200	6606.451613	12412.12121	37236.36364
Exp 2	3200	5851.428571	11702.85714	34133.33333	3303.225806	6400	13653.33333	37236.36364
Exp 3	3103.030303	6023.529412	11377.77778	34133.33333	3303.225806	6826.666667	12800	39009.52381
Average	3238.787879	5908.795518	11594.49735	34628.01932	3268.817204	6611.039427	12955.15152	37827.41703
Standard Dev.	129.6157584	81.12911448	153.2438799	699.5916362	48.66111161	174.2161439	518.4630336	835.8757202



## Latency:

Test Results for 8B								
	Sequential R+W				Random R+W			
	1 Thread	2Thread	4Thread	8Thread	1 Thread	2Thread	4Thread	8Thread
Exp 1	7.82E-07	4.14E-07	2.03E-07	1.00E-07	4.71E-06	3.54E-05	1.89E-05	9.74E-06
Exp 2	7.26E-07	4.05E-07	1.98E-07	1.02E-07	4.45E-06	3.60E-05	1.67E-05	5.34E-06
Exp 3	7.54E-07	4.05E-07	1.98E-07	1.02E-07	4.49E-06	3.33E-05	1.79E-05	8.07E-06
Average	7.54E-07	4.08E-07	1.99E-07	1.02E-07	4.55E-06	3.49E-05	1.78E-05	7.72E-06
Standard Dev.	2.28E-08	4.39E-09	2.20E-09	1.10E-09	1.15E-07	1.19E-06	8.87E-07	1.81E-06

## STREAM BENCHMARK:

Below is the screenshot of STREAM benchmarking showing the results of Memory Speeds.

The result below shows the memory speed having a total of 26.54GB/s



```
[cc@pa1-shushupti-prasanna ~]$ gcc -o stream stream.c
[cc@pa1-shushupti-prasanna ~]$ ./stream
-----
STREAM version $Revision: 5.10 $
-----
This system uses 8 bytes per array element.
-----
Array size = 10000000 (elements), Offset = 0 (elements)
Memory per array = 76.3 MiB (= 0.1 GiB).
Total memory required = 228.9 MiB (= 0.2 GiB).
Each kernel will be executed 10 times.
The *best* time for each kernel (excluding the first iteration)
will be used to compute the reported bandwidth.
-----
Your clock granularity/precision appears to be 1 microseconds.
Each test below will take on the order of 32410 microseconds.
(= 32410 clock ticks)
Increase the size of the arrays if this shows that
you are not getting at least 20 clock ticks per test.
-----
WARNING -- The above is only a rough guideline.
For best results, please be sure you know the
precision of your system timer.
-----
Function      Best Rate MB/s  Avg time     Min time     Max time
Copy:         5843.3    0.029205     0.027382     0.031202
Scale:        5659.3    0.029841     0.028272     0.031645
Add:          7962.9    0.030738     0.030140     0.032213
Triad:        7673.9    0.032267     0.031275     0.033335
-----
Solution Validates: avg error less than 1.000000e-13 on all three arrays
-----
```

## Disk Benchmarking

In this experiment, we are performing Disk Benchmarking across 1,2,4 and 8 threads to measure the efficiency of concurrency. Thread operations are done on different block size of data - 8B,8KB,8MB and 80MB and the operations are Sequential read+write, sequential read access and random read access. We have taken a 1GB file to perform the read/write operations on the disk.

Below are the graphs showing the experimental results of the program we written on C:




Throughput are plotted for 8 KB,8 MB and 80 MB cases.



From the above graphs, it is seen that the throughput increases as we increase the number of threads.

## Latency:

Latency is plotted considering the 8B case



From the above graphs it can be seen that the latency keeps decreasing as we increase the number of threads. However for random read access, they seem to increase at use of 2 threads but the latency decreased as we increased the threads to 4 and 8.

## Test Results:

	Sequential Read+Write - 8KB			
	1 Thread	2Thread	4Thread	8Thread
Exp 1	1101.075269	1197.660819	3011.764706	5610.958904
Exp 2	975.238095	1226.347305	2712.582781	5572.789116
Exp 3	1101.075269	1044.897959	2748.993289	5497.986577
Average	1059.129544	1156.302028	2824.446925	5560.578199
Standard Dev.	59.32021271	79.64035203	133.2851442	46.92203844

Test Results for 8KB								
	Sequential Read				Random Read			
	1 Thread	2Thread	4Thread	8Thread	1 Thread	2Thread	4Thread	8Thread
Exp 1	4096	10240	22755.55556	35617.3913	5120	6826.666667	16384	29257.14286
Exp 2	4654.545455	8904.347826	19504.76191	37236.36364	6826.666667	9309.090909	17808.69565	35617.3913
Exp 3	4096	8192	17808.69565	45511.11111	7314.285714	8533.333333	17808.69565	35617.3913
Average	4282.181818	9112.115942	20023.00437	39454.95535	6420.31746	8223.030303	17333.7971	33497.30849
Standard Dev.	263.3008526	848.9019152	2052.524816	4333.05395	940.7664276	1036.926068	671.6079702	2998.249868

	Sequential Read+Write - 8MB			
	1 Thread	2Thread	4Thread	8Thread
Exp 1	1312.820513	2226.086957	4602.247191	8110.891089
Exp 2	1296.202532	2250.549451	4266.666667	8192
Exp 3	1296.202532	2409.411765	4357.446809	8192
Average	1301.741859	2295.349391	4408.786889	8164.963696
Standard Dev.	7.83379137	81.27021707	141.7284372	38.23510732

Test Results for 8MB								
	Sequential Read				Random Read			
	1 Thread	2Thread	4Thread	8Thread	1 Thread	2Thread	4Thread	8Thread
Exp 1	4266.666667	7062.068966	16384	30340.74074	7876.923077	14628.57143	24094.11765	51200
Exp 2	4654.545455	7062.068966	15170.37037	30340.74074	10240	14628.57143	29257.14286	54613.33333
Exp 3	4654.545455	9309.090909	17066.66667	30340.74074	8533.333333	17066.66667	27306.66667	51200
Average	4525.252526	7811.07628	16207.01235	30340.74074	8883.418803	15441.26984	26885.97573	52337.77778
Standard Dev.	182.8478142	1059.256302	784.2101701	3.64E-12	995.9762425	1149.329118	2128.683941	1609.060763

	Sequential Read+Write - 80MB			
	1 Thread	2Thread	4Thread	8Thread
Exp 1	1264.197531	2068.686869	4452.173913	11221.91781
Exp 2	1204.705882	2202.150538	4602.247191	11538.02817
Exp 3	1066.666667	2007.843137	4179.591837	11377.77778
Average	1178.52336	2092.893515	4411.337647	11379.24125
Standard Dev.	82.73956692	81.15135157	174.947778	129.055663

Test Results for 80MB								
	Sequential Read				Random Read			
	1 Thread	2Thread	4Thread	8Thread	1 Thread	2Thread	4Thread	8Thread
Exp 1	3531.034483	6826.666667	12800	31507.69231	4096	7585.185185	12412.12121	22140.54054
Exp 2	3531.034483	6400	13212.90323	31507.69231	3938.461538	7585.185185	11702.85714	18204.44444
Exp 3	3531.034483	6606.451613	12412.12121	31507.69231	4266.666667	7876.923077	14628.57143	32768
Average	3531.034483	6611.039427	12808.34148	31507.69231	4100.376068	7682.431149	12914.51659	24370.99499
Standard Dev.	0	174.2161439	326.9710957	0	134.0249085	137.5265612	1246.127834	6151.177824

## Latency:

	Sequential Read+Write - 8B			
	1 Thread	2Thread	4Thread	8Thread
Exp 1	1.19E-06	5.05E-07	2.47E-07	1.25E-07
Exp 2	9.48E-07	4.84E-07	2.44E-07	1.27E-07
Exp 3	8.64E-07	4.90E-07	2.42E-07	1.25E-07
Average	1.00E-06	4.93E-07	2.44E-07	1.25E-07
Standard Dev.	1.38E-07	9.03E-09	2.27E-09	8.09E-10

Test Results for 8B								
	Sequential Read				Random Read			
	1 Thread	2Thread	4Thread	8Thread	1 Thread	2Thread	4Thread	8Thread
Exp 1	3.88E-07	2.14E-07	1.04E-07	5.25E-08	3.24E-06	3.97E-06	2.49E-06	1.07E-06
Exp 2	4.40E-07	2.17E-07	1.06E-07	5.28E-08	3.25E-06	4.39E-06	2.04E-06	9.91E-07
Exp 3	3.86E-07	2.10E-07	1.02E-07	5.09E-08	3.57E-06	4.26E-06	2.29E-06	1.02E-06
Average	4.05E-07	2.14E-07	1.04E-07	5.21E-08	3.35E-06	4.21E-06	2.27E-06	1.03E-06
Standard Dev.	2.50E-08	2.85E-09	1.71E-09	8.32E-10	1.54E-07	1.75E-07	1.83E-07	3.47E-08

## IoZone Benchmarking:

The following screenshot shows the performance metrics taken from IoZone benchmarking ran on the system. The speeds are measured in terms of Kilobytes. The first image shows the speeds of all values which gets run when executing the benchmark. The second image is for a specific case where we have given the filesize.

Run began: Sun Oct 8 21:26:51 2017

Auto Mode

Command line used: ./iozone -a

Output is in kBytes/sec

Time Resolution = 0.000001 seconds.

Processor cache size set to 1024 kBytes.

Processor cache line size set to 32 bytes.

File stride size set to 17 \* record size.

kB	reclen	write	rewrite	read	reread	random read	random write	bkwd read	record rewrite	stride read	fwrite	frewrite	fread	freread
64	4	673098	1232453	3791156	4018152	2601873	1330167	1780006	1638743	2067979	1330167	1421755	2358717	3203069
64	8	889431	1599680	5860307	5389653	4274062	1722886	2203800	2133730	2892445	1562436	1599680	2772930	4274062
64	16	811459	1119387	5735102	7100397	4988978	1991276	2278628	1991276	2772930	1562436	1828508	2892445	4564786
64	32	901377	1734015	5860307	5860307	4274062	1991276	2278628	1232453	3057153	1734015	2203800	3165299	4897948
64	64	952555	1722886	6421025	7100397	4564786	1991276	2358717	1991276	2801873	1879725	1933893	3363612	5283570
128	4	853794	1471662	2132084	4135958	3199360	1487977	2202044	1905110	2464907	1598754	1504659	2969325	3560017
128	8	962469	1802755	5325799	5603747	4407601	1975201	2920861	2511022	4267461	1852520	1878447	4012317	4717434
128	16	809997	1618028	5764891	6406138	5122535	2326073	3468030	2727923	4407601	2202044	2238774	4135958	5325799
128	32	1185654	2511022	5545660	6045455	5325799	2511022	3657016	2905056	4759253	2727923	2286447	4104338	2606629
128	64	955616	2166499	5847904	6727225	5325799	2558895	3468030	2784517	2367096	1878447	2969325	3895854	3380677
128	128	1022989	2166499	6727225	5847904	2673584	2248149	2621367	4012317	2409592	1939522	4012317	5122535	
256	4	886067	2327564	4261170	4496299	3454704	1292410	2872459	2114473	3009318	1752173	1610276	3410608	3756894
256	8	873809	2015259	5320671	5938650	4840911	2330140	3605511	1965448	3823789	1802167	1766587	3705040	5569035
256	16	1157291	2371308	6249745	4734192	5569035	2639446	3935921	2819657	3605511	2210228	2210228	4422226	5347168
256	32	1195962	1652404	3935921	5569035	5117791	2665657	3935921	3078337	4572895	2413957	1740810	4553502	5242734
256	64	1219045	2419396	5687020	5142301	3823789	2849590	1487577	2975955	5117791	2371308	2719671	4652146	5320671
256	128	1206714	2305128	3879045	5810112	5455847	3087188	4117018	3368013	4907284	2019049	3454704	4572895	5320671
256	256	870974	2392442	5455847	5569035	5320671	2783115	2783115	2719671	4332998	2783115	2726577	4907284	5117791
512	4	932550	1662389	4030519	4163357	3239989	1508098	2942541	2178404	3029721	1594501	1547395	2975154	2497636
512	8	1007827	2073249	5115422	5624545	4784885	2275345	3510074	2584820	4198596	2132967	1712772	4871723	5331314
512	16	1221684	2317060	5822804	5951911	4447925	2613128	4494470	3459188	4742616	2285029	1848401	5227492	5624545
512	32	1171891	2073249	3878421	5624545	5227492	2845061	4871723	2910635	5115422	2438090	1940253	4131319	5384786
512	64	1078159	2544997	5822804	5951911	5744919	3029721	3970896	3481620	5398323	2668321	2651850	5019764	5624545
512	128	1248678	2708713	4532414	3877684	5509113	2497638	3877684	3556580	5331314	2285029	2926501	4660280	6414119
512	256	1289930	2753870	5760329	5886650	4393943	3239989	4973263	3264616	5278892	1821745	3762197	4784885	5278892
512	512	1017858	2535961	3684733	5684095	5384786	2958758	4742616	3182373	4571003	3177664	2975154	5019764	5384786
1024	4	1021928	1638171	4029784	4249053	2737730	1646334	2572135	2160657	3020783	1837965	1525848	3835457	3941039
1024	8	1120288	2173780	6128613	5593820	3880507	2333199	4410497	2985091	4550690	2031864	1954207	5200941	5356618
1024	16	1311117	2305644	5821271	6059442	5991815	2869422	5789881	3778101	4876180	2491561	2480051	3738636	4079544
1024	32	1311117	2573677	5623115	4995276	5479632	2737730	5417427	3970183	4949226	2708381	2708381	5789881	4574926
1024	64	1343938	2432298	5821271	5917516	7862087	2976816	5330028	3906759	4920874	2358827	2653170	5925680	5444898
1024	128	1399100	2375790	4232305	5853003	4654248	3210456	4130547	5660167	2716948	2716948	2715230	4300102	5623115
1024	256	1374474	2370545	4762630	5472650	5479632	3103735	4995276	3579719	5336651	2708381	2985091	5194651	4898425
1024	512	1238168	2791104	5500686	5564829	5536137	3083680	5048117	3379719	5226256	2102484	3880507	3631168	5390231
1024	1024	1192117	2552265	4805258	5593820	5251818	3092561	5096034	3130886	5652717	3314514	2730767	4926518	4634161
2048	4	941176	1823453	3962502	3901313	3121562	1859773	2920945	2397544	3324528	1725646	1344529	3938881	3730203
2048	8	1265122	2098276	4686767	5640860	4284671	2484934	4068930	3271351	5041616	2387548	2164908	5349298	4911886
2048	16	1371798	2412357	5904523	6058611	4490751	2824867	5212953	3374151	5719737	2479196	2260624	5750369	5916724
2048	32	1145965	2400894	5578583	5107566	5517666	2805512	4934459	3764533	5251194	2522144	2423929	4807425	5596757
2048	64	1261036	1832009	4983122	5920802	5689430	3209024	5593112	3374151	5481535	2208455	1933883	5389574	6464402
2048	128	1163663	2578935	6058611	6243566	5578583	3374151	4707314	4094723	5251194	2678988	2470639	4751580	5644567
2048	256	1380173	2060032	4889516	5478461	5489457	3000267	4623699	3483621	5447681	2540795	2531808	5567735	5852228
2048	512	1444219	1890887	5492967	5535444	5032754	2994254	4807425	3414387	4643695	2920945	2172756	4601407	5475461
2048	1024	1183707	2045804	4831762	5640860	5611381	2981782	4772701	3418463	5447681	2235910	3682233	4720248	5582206
2048	2048	1204792	2619037	5159722	5769681	6189580	3112514	3885431	2933915	5403134	2999482	2884652	5293261	5156624

[cc@pal-shushupti-prasanna current]\$. ./iozone -g# -s 8192

Iozone: Performance Test of File I/O

Version \$Revision: 3.471 \$

Compiled for 64 bit mode.

Build: linux-ia64

Contributors:William Norcott, Don Capps, Isom Crawford, Kirby Collins  
Al Slater, Scott Rhine, Mike Wisner, Ken Goss  
Steve Landherr, Brad Smith, Mark Kelly, Dr. Alain CYR,  
Randy Dunlap, Mark Montague, Dan Million, Gavin Brebner,  
Jean-Marc Zucconi, Jeff Blomberg, Benny Halevy, Dave Boone,  
Erik Habbinga, Kris Strecker, Walter Wong, Joshua Root,  
Fabrice Bacchella, Zhenghua Xue, Qin Li, Darren Sawyer,  
Vangel Bojaxhi, Ben England, Vikentsi Lapa,  
Alexey Skidanov.

Run began: Tue Oct 10 02:04:32 2017

Using maximum file size of 4 kilobytes.

File size set to 8192 kB

Command line used: ./iozone -g# -s 8192

Output is in kBytes/sec

Time Resolution = 0.000001 seconds.

Processor cache size set to 1024 kBytes.

Processor cache line size set to 32 bytes.

File stride size set to 17 \* record size.

kB	reclen	write	rewrite	read	reread	random read	random write	bkwd read	record rewrite	stride read	fwrite	frewrite	fread	freread
8192	4	653838	1373576	4222755	4576535	3497955	1858005	2583282	2383476	3415886	1984069	1460016	3291223	4011735


iozone test complete.

Performance tested for 8MB(8192 KB) file size :

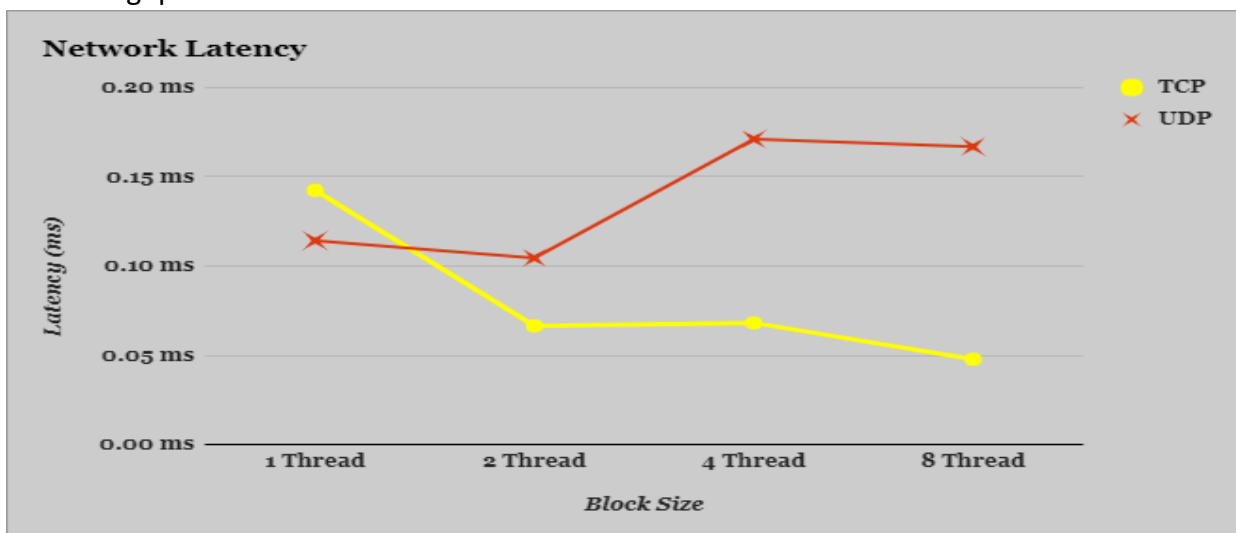
From the above we can note down the speed performances, sequential write -0.62 Gb/s, sequential read - 4.03 Gb/s and random read -3.33 Gb/s

## Network Benchmarking

- Network Benchmark evaluation is performed for both TCP and UDP using fixed scaling.
- It includes the TCP protocol stack, UDP, fixed packet/buffer size (64KB), and varying the concurrency (1 thread, 2 threads, 4 threads, 8 threads);
- Size of data is 1GB.
- The throughput and latency for each of them is shown in the below graph:



- Here we can see that throughput is increasing as the threads are increasing for TCP connection. But, for UDP, throughput is decreasing. As UDP being Connectionless protocol is slower as compared to TCP. There might be lost of packets in case of UDP. Big Data size is also a reason for decreasing throughput.



- Here we can see that latency is decreasing as the threads are increasing for TCP connection. But, for UDP its increasing. Latency gives us the idea about the continuous sending of data, so it should be decreasing. UDP is a connectionless protocol, hence its taking time to send the data.



- **Throughput and Latency for network:**

Network Throughput	1 Thread	2 Thread	4 Thread	8 Thread
TCP	45.99979	98.34523	95.98722	136.7811
UDP	57.02899	62.63849	38.31422	39.26602

Network Latency	1 Thread	2 Thread	4 Thread	8 Thread
TCP	0.14247	0.066639	0.068276	0.047913
UDP	0.114258	0.104626	0.171049	0.166903

- **Mean Throughput and Latency for three experiments:**

Mean Value	Thread 1	Thread 2	Thread 4	Thread 8
TCP	43.48793	99.47256	94.94335	134.8511
UDP	57.2036	73.71618	36.71421	43.46394

Mean Value	Thread 1	Thread 2	Thread 4	Thread 8
TCP	0.095307	0.049078	0.06776	0.035685
UDP	0.157589	0.101295	0.191048	0.267643

- **Standard Deviation of Throughput and Latency for three experiments:**

SD	Thread 1	Thread 2	Thread 4	Thread 8
TCP	4.481194	4.165044	5.60816	3.342858
UDP	6.271823	13.59424	3.23209	3.661663


SD	Thread 1	Thread 2	Thread 4	Thread 8
TCP	0.051063	0.035739	0.006243	0.014802
UDP	0.075052	0.005768	0.034639	0.174487

- IPERF Benchmarking for TCP and UDP:

```
Client connecting to 52.26.161.78, TCP port 5001
TCP window size: 325 KByte (default)
-----
[ 3] local 172.31.25.167 port 46111 connected with 52.26.161.78 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.0 sec    128 KBytes  18.7 Gbits/sec
ubuntu@ip-172-31-25-167:~$ iperf -c 52.26.161.78 -n 64
-----
Client connecting to 52.26.161.78, TCP port 5001
TCP window size: 325 KByte (default)
-----
[ 3] local 172.31.25.167 port 46112 connected with 52.26.161.78 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.0 sec    128 KBytes  18.7 Gbits/sec
ubuntu@ip-172-31-25-167:~$ iperf -c 52.26.161.78 -n 1024
-----
Client connecting to 52.26.161.78, TCP port 5001
TCP window size: 325 KByte (default)
-----
[ 3] local 172.31.25.167 port 46113 connected with 52.26.161.78 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.0 sec    128 KBytes  15.7 Gbits/sec
ubuntu@ip-172-31-25-167:~$ iperf -c 52.26.161.78 -n 65536
-----
Client connecting to 52.26.161.78, TCP port 5001
TCP window size: 325 KByte (default)
-----
[ 3] local 172.31.25.167 port 46114 connected with 52.26.161.78 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.0 sec    128 KBytes  16.4 Gbits/sec
```


```
Client connecting to 52.26.161.78, UDP port 5001
Sending 1470 byte datagrams
UDP buffer size: 208 KByte (default)
-----
[ 3] local 172.31.25.167 port 50114 connected with 52.26.161.78 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.0 sec    1.44 KBytes  1.04 Mbits/sec
[ 3] Sent 1 datagrams
[ 3] Server Report:
[ 3] 0.0- 0.0 sec    1.44 KBytes  1.04 Mbits/sec    0.000 ms    0/    1 (0%)
ubuntu@ip-172-31-25-167:~$ iperf -c 52.26.161.78 -u -n 1024
-----
Client connecting to 52.26.161.78, UDP port 5001
Sending 1470 byte datagrams
UDP buffer size: 208 KByte (default)
-----
[ 3] local 172.31.25.167 port 56925 connected with 52.26.161.78 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.0 sec    1.44 KBytes  1.04 Mbits/sec
[ 3] Sent 1 datagrams
[ 3] Server Report:
[ 3] 0.0- 0.0 sec    1.44 KBytes  1.05 Mbits/sec    0.000 ms    0/    1 (0%)
ubuntu@ip-172-31-25-167:~$ iperf -c 52.26.161.78 -u -n 65536
-----
Client connecting to 52.26.161.78, UDP port 5001
Sending 1470 byte datagrams
UDP buffer size: 208 KByte (default)
-----
[ 3] local 172.31.25.167 port 55050 connected with 52.26.161.78 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0- 0.5 sec    64.6 KBytes  1.05 Mbits/sec
[ 3] Sent 45 datagrams
[ 3] Server Report:
[ 3] 0.0- 0.5 sec    64.6 KBytes  1.05 Mbits/sec    0.041 ms    0/   45 (0%)
```

- Network Benchmarking on two different nodes:



Network Throughput	1 Thread	2 Thread	4 Thread	8 Thread
TCP	108.2749328	229.7124246	109.3656836	276.0400275
UDP	111.2610902	108.8259378	57.06838789	59.76927625

- We can see that throughput is increasing as the threads are increasing for TCP connection. But, for UDP, throughput is decreasing.



Network Latency	1 Thread	2 Thread	4 Thread	8 Thread
TCP	0.06052739844	0.02852958438	0.05992373281	0.02374148438
UDP	0.05890289219	0.0602209375	0.1148376578	0.1096483078

- Here we can see that latency is decreasing as the threads are increasing for TCP connection. But, for UDP its increasing. Latency gives us the idea about the continuous sending of data, so it should be decreasing.